

```
#Importing the libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
sns.set()
import warnings
warnings.filterwarnings('ignore')

#loading the dataset
df = pd.read_csv('Sleep_health_and_lifestyle_dataset.csv')
df.head()
```

	Person ID	Gender	Age	Occupation	Sleep Duration \
0	1	Male	27	Software Engineer	6.1
1	2	Male	28	Doctor	6.2
2	3	Male	28	Doctor	6.2
3	4	Male	28	Sales Representative	5.9
4	5	Male	28	Sales Representative	5.9

	Quality of Sleep Category \	Physical Activity Level	Stress Level	BMI
0		6	42	6
1	Overweight	6	60	8
2	Normal	6	60	8
3	Normal	4	30	8
4	Obese	4	30	8

	Blood Pressure	Heart Rate	Daily Steps	Sleep Disorder
0	126/83	77	4200	NaN
1	125/80	75	10000	NaN
2	125/80	75	10000	NaN
3	140/90	85	3000	Sleep Apnea
4	140/90	85	3000	Sleep Apnea

## Some Numerical Information about the Data

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Person ID                            374 non-null    int64
```

1	Gender	374	non-null	object
2	Age	374	non-null	int64
3	Occupation	374	non-null	object
4	Sleep Duration	374	non-null	float64
5	Quality of Sleep	374	non-null	int64
6	Physical Activity Level	374	non-null	int64
7	Stress Level	374	non-null	int64
8	BMI Category	374	non-null	object
9	Blood Pressure	374	non-null	object
10	Heart Rate	374	non-null	int64
11	Daily Steps	374	non-null	int64
12	Sleep Disorder	155	non-null	object

dtypes: float64(1), int64(7), object(5)  
memory usage: 38.1+ KB

```
df.nunique()
```

Person ID	374
Gender	2
Age	31
Occupation	11
Sleep Duration	27
Quality of Sleep	6
Physical Activity Level	16
Stress Level	6
BMI Category	4
Blood Pressure	25
Heart Rate	19
Daily Steps	20
Sleep Disorder	2

dtype: int64

## Data Cleaning

```
# Fill NaN values with 'Healthy Sleep'
df['Sleep Disorder'].fillna('Healthy Sleep', inplace=True)
df['Sleep Disorder'].value_counts()

Sleep Disorder
Healthy Sleep    219
Sleep Apnea      78
Insomnia         77
Name: count, dtype: int64

# Create Age Group
df['Age'] = pd.cut(df['Age'], bins=[26, 30, 35, 40, 45, 50, 55, 60],
labels=['26-30', '30-35', '35-40', '40-45', '45-50', '50-55', '55-60'])
df['Age'].value_counts()
```

```
Age
40-45    99
35-40    71
30-35    62
50-55    43
45-50    34
55-60    33
26-30    32
Name: count, dtype: int64
```

```
company_dic = df['Occupation'].value_counts().head(6)
```

```
def occup(x, dic):
    if x in dic.keys():
        return x
    else :
        return 'Others'
```

```
df['Occupation'] = df['Occupation'].apply(lambda x : occup(x,
company_dic))
df['Occupation'].value_counts()
```

```
Occupation
Nurse        73
Doctor       71
Engineer     63
Lawyer       47
Others       43
Teacher      40
Accountant   37
Name: count, dtype: int64
```

```
# Reduce unique values of BMI Category
```

```
df['BMI Category'] = df['BMI Category'].apply(lambda x : 'Normal' if x
== 'Normal Weight' else x)
df = df[df['BMI Category'] != 'Obese']
df['BMI Category'].value_counts()
```

```
BMI Category
Normal      216
Overweight  148
Name: count, dtype: int64
```

```
#splitting the blood pressure into two columns
```

```
df['Systolic BP'] = df['Blood Pressure'].apply(lambda x: x.split('/')[
0])
df['Diastolic BP'] = df['Blood Pressure'].apply(lambda x: x.split('/')[
1])
```

```
#dropping the blood pressure column
```

```
df.drop('Blood Pressure', axis=1, inplace=True)
```

# Data Visualization

```
# Define list of Continuous columns Names
continuous = ['Sleep Duration', 'Quality of Sleep', 'Physical Activity
Level', 'Stress Level', 'Heart Rate', 'Daily Steps', 'Systolic BP',
'Diastolic BP']

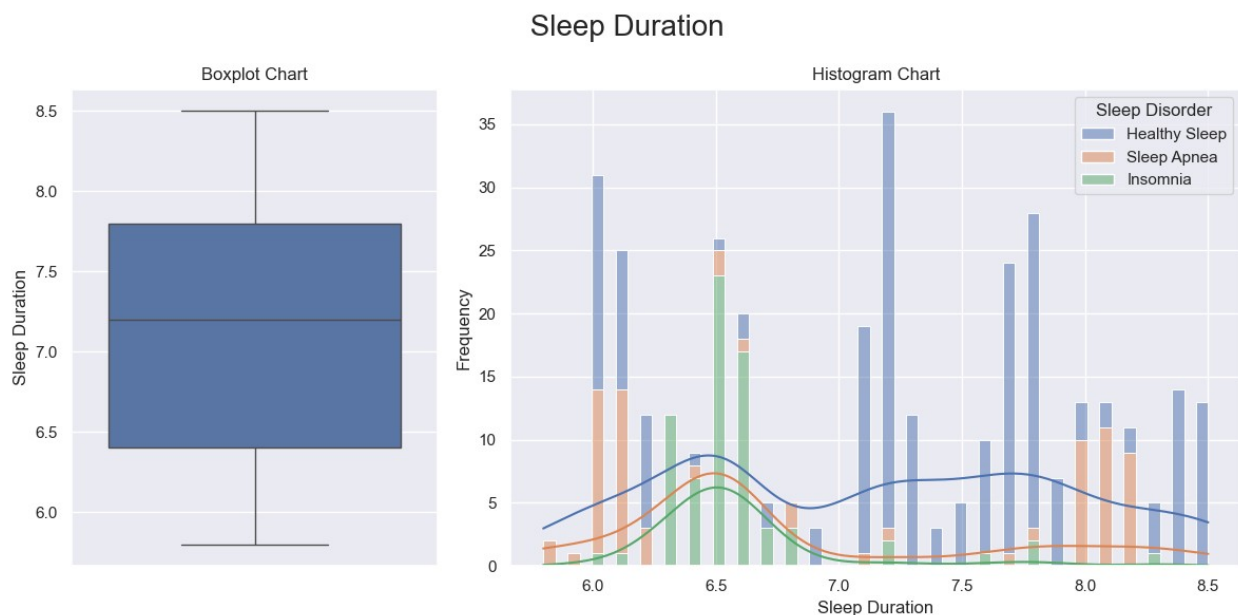
# Distribution of Categorical Features
def plot_continuous_distribution(df, column, hue):

    width_ratios = [2, 4]
    gridspec_kw = {'width_ratios':width_ratios}
    fig, ax = plt.subplots(1, 2, figsize=(12, 6), gridspec_kw =
gridspec_kw)
    fig.suptitle(f' {column} ', fontsize=20)

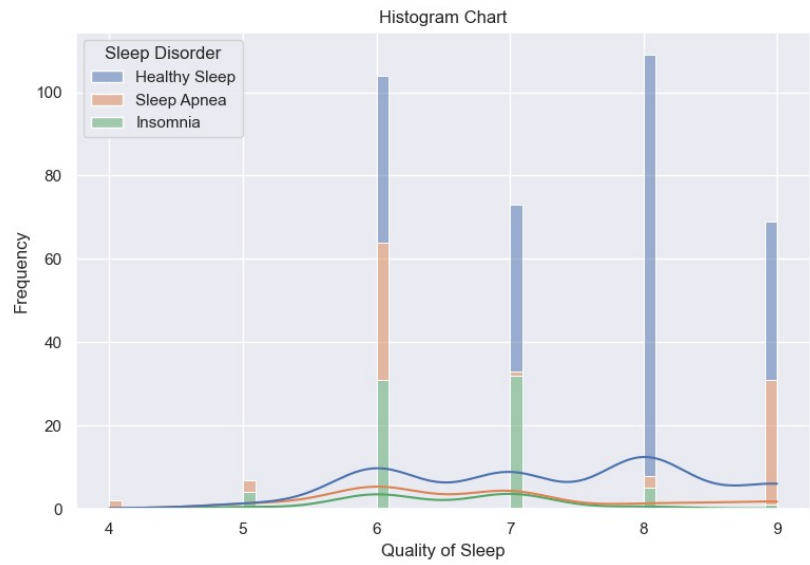

    sns.boxplot(df[column], ax=ax[0])
    ax[0].set_title('Boxplot Chart')
    ax[0].set_ylabel(column)

    sns.histplot(x = df[column], kde=True, ax=ax[1], hue=df[hue],
multiple = 'stack', bins=55)
    ax[1].set_title('Histogram Chart')
    ax[1].set_ylabel('Frequency')
    ax[1].set_xlabel(column)



    plt.tight_layout()
    plt.show()
for conti in continuous :
    plot_continuous_distribution(df, conti, 'Sleep Disorder')
```



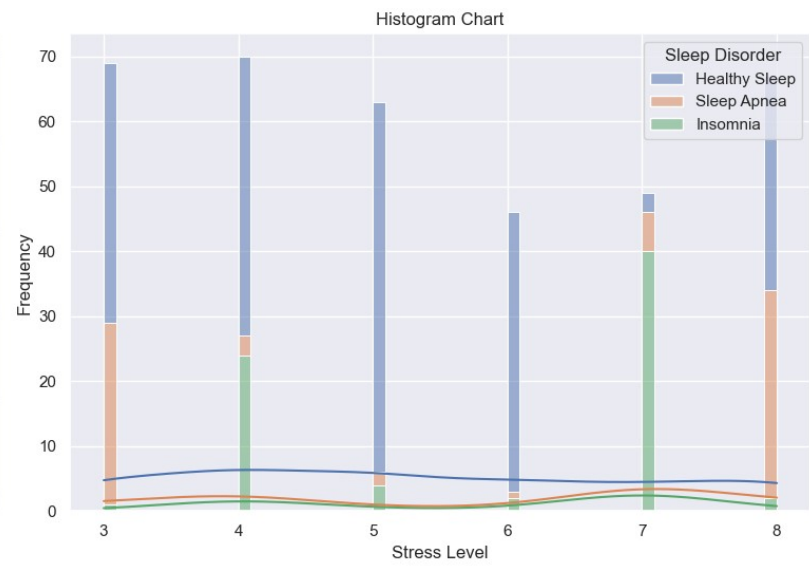

## Quality of Sleep



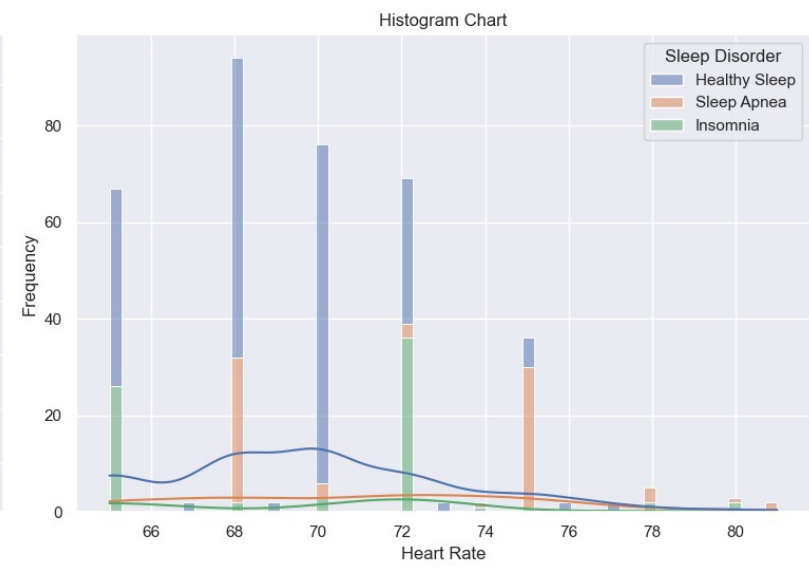

## Physical Activity Level





## Stress Level



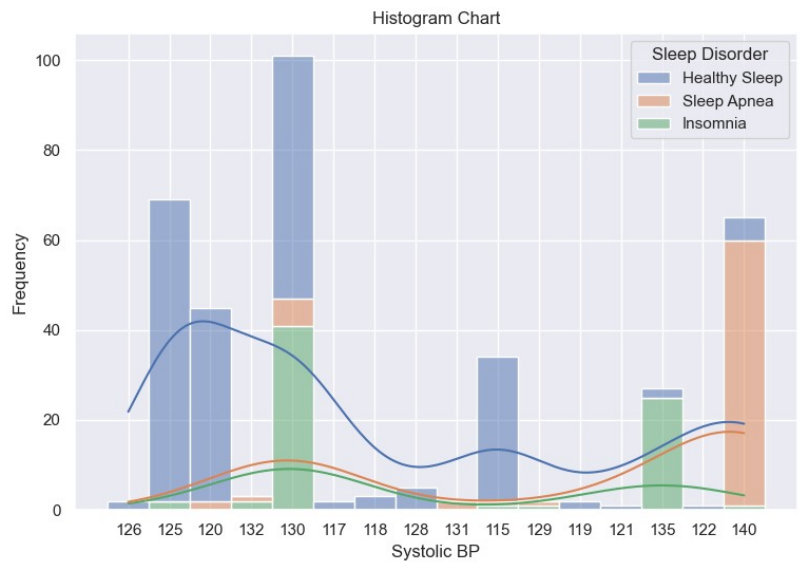

## Heart Rate



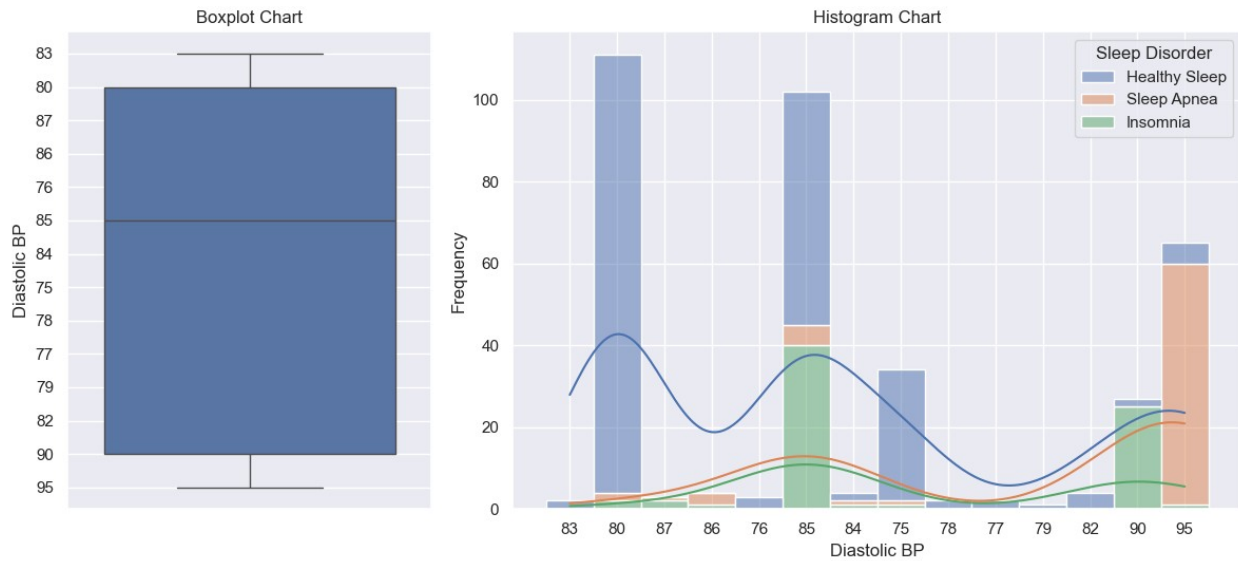
## Daily Steps



## Systolic BP



## Diastolic BP


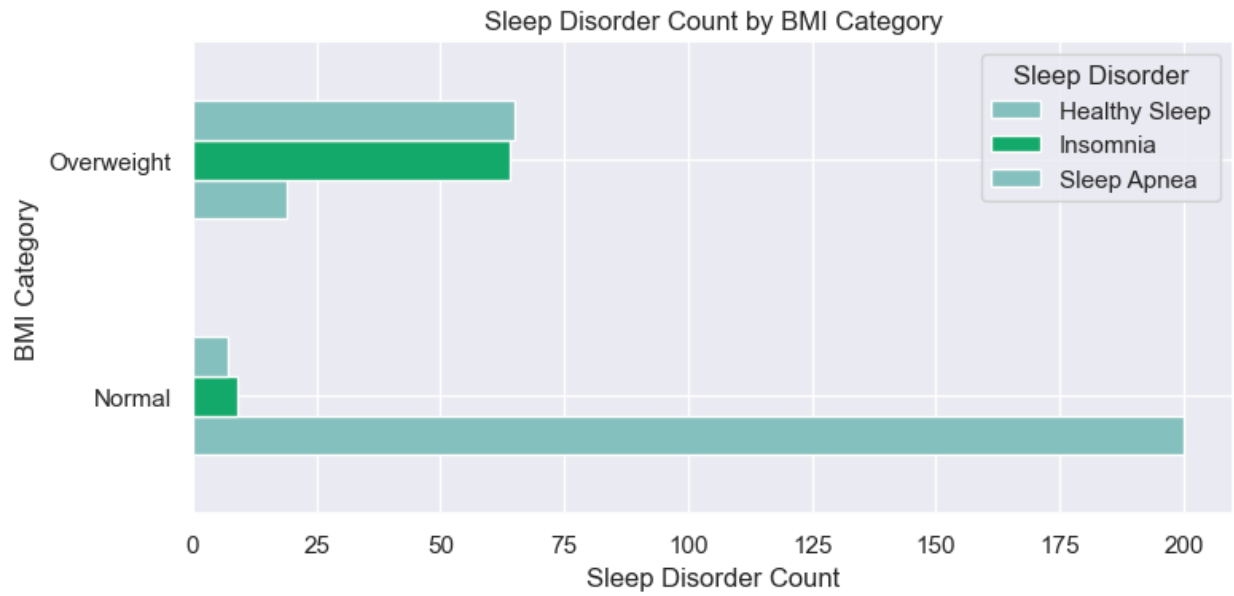


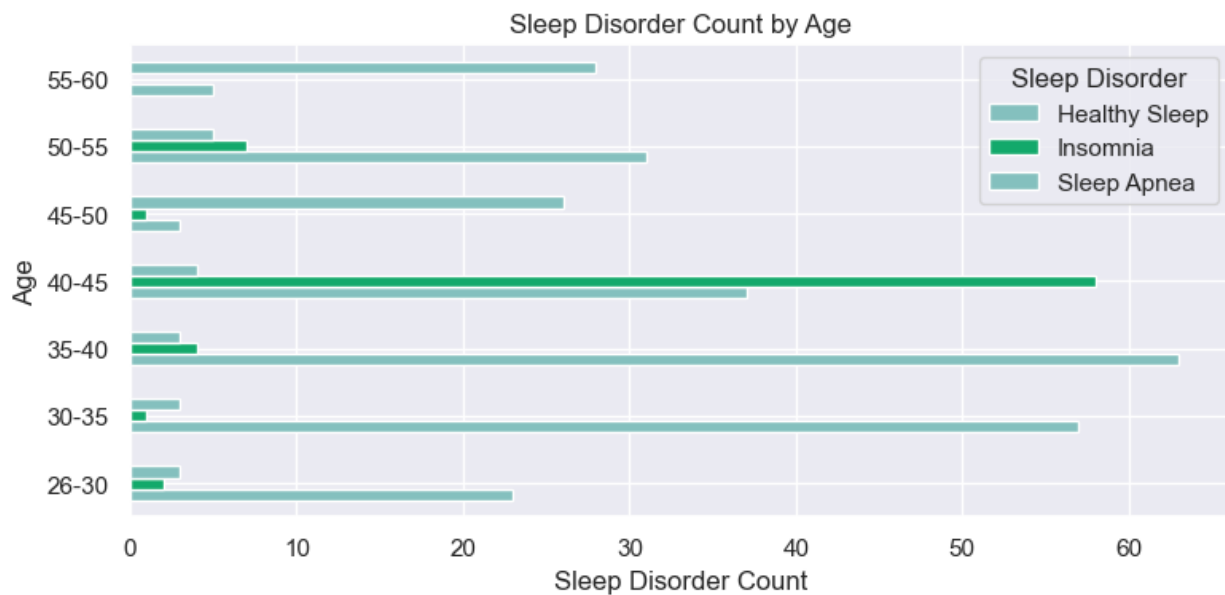

```
# Define a Function for Barh Plot
def bar_plot(x, y, df):
    barh = df.groupby([x, y]).size().unstack()
    barh.plot(kind='barh', color = ['#84c0be', '#13a96b'],
figsize=(8,4))
    plt.title(f'{y} Count by {x}')
    plt.xlabel(f'{y} Count')
    plt.ylabel(x)

    plt.tight_layout()
    plt.show()

bar_plot('BMI Category', 'Sleep Disorder', df)
bar_plot('Quality of Sleep', 'Sleep Disorder', df)
bar_plot('Occupation', 'Sleep Disorder', df)
bar_plot('Age', 'Sleep Disorder', df)
```







## Data Preprocessing

```
from sklearn.preprocessing import StandardScaler, LabelEncoder

# Initialize StandardScaler
stc = StandardScaler()
# Initialize Label Encoder
le = LabelEncoder()

stc_cols = ['Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'systolic_bp', 'diastolic_bp', 'Heart Rate', 'Daily Steps']
dum_cols= ['Age', 'Occupation', 'BMI Category']
```

```

le_cols = ['Gender', 'Sleep Disorder']

# Apply Standard Scaler to the selected columns
df[stc_cols] = stc.fit_transform(df[stc_cols])

# Apply Label Encoder to the selected columns
for col in le_cols :
    df[col] = le.fit_transform(df[col])

# Apply get_dummies to the selected columns
df = pd.get_dummies(df, columns=dum_cols)

```

## Training and Evaluating Different Models

```

from sklearn.model_selection import train_test_split

x = df.drop(['Sleep Disorder', 'Person ID'], axis=1)
y = df['Sleep Disorder'] # Target Variable

x_train, x_test, y_train, y_test = train_test_split(x, y,
test_size=0.15, random_state=12)

#Importing the Libraries
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import accuracy_score
from xgboost import XGBClassifier

# List of Models to Try
models = [
    ('Gradient Boosting', GradientBoostingClassifier()),
    ('K-Nearest Neighbors', KNeighborsClassifier()),
    ('Random Forest', RandomForestClassifier()),
    ('Decision Tree', DecisionTreeClassifier()),
    ('XGB Classifier', XGBClassifier())
]

# Train and evaluate each model
for name, model in models:
    model.fit(x_train, y_train)
    y_pred = model.predict(x_test)
    print(f'Training accuracy: {name}', model.score(x_train, y_train))
    print(f'Test accuracy: {name}', accuracy_score(y_test, y_pred))
    print()

```

Training accuracy: Gradient Boosting 0.9255663430420712  
Test accuracy: Gradient Boosting 0.9454545454545454

Training accuracy: K-Nearest Neighbors 0.9061488673139159  
Test accuracy: K-Nearest Neighbors 0.9090909090909091

Training accuracy: Random Forest 0.9255663430420712  
Test accuracy: Random Forest 0.9454545454545454

Training accuracy: Decision Tree 0.9255663430420712  
Test accuracy: Decision Tree 0.9454545454545454

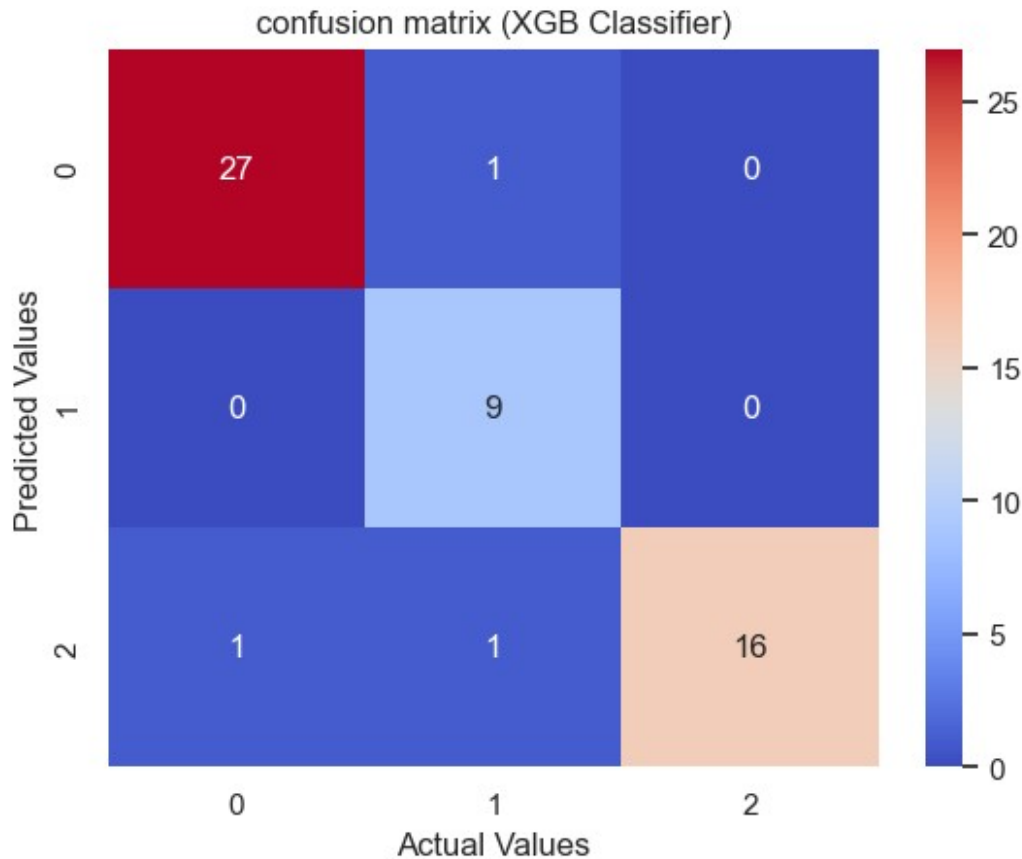
Training accuracy: XGB Classifier 0.9255663430420712  
Test accuracy: XGB Classifier 0.9454545454545454

```
xgb = XGBClassifier()  
xgb.fit(x_train, y_train)  
xgb_pred = xgb.predict(x_test)
```

```
accuracy = accuracy_score(y_test, xgb_pred)  
print(f'R-squared (XGB Classifier): {round(accuracy, 3)}')
```

R-squared (XGB Classifier): 0.945

```
from sklearn.metrics import confusion_matrix, classification_report  
sns.heatmap(confusion_matrix(y_test, xgb_pred), annot=True, cmap =  
'coolwarm', fmt='.0f')  
plt.ylabel('Predicted Values')  
plt.xlabel('Actual Values')  
plt.title('confusion matrix (XGB Classifier)')  
plt.show()
```



```
# Visualize Classification report for XGB Classifier
print(classification_report(y_test,xgb_pred))
```

	precision	recall	f1-score	support
0	0.96	0.96	0.96	28
1	0.82	1.00	0.90	9
2	1.00	0.89	0.94	18
accuracy			0.95	55
macro avg	0.93	0.95	0.94	55
weighted avg	0.95	0.95	0.95	55

## Summary and Conclusion

In this project, I focused on predicting sleep disorders using a variety of data preprocessing and machine learning techniques. The steps and methodologies employed are as follows:

1. Data Cleaning and Preprocessing:
  - Handling Missing Values: Missing values were filled in a way that preserved the distribution of the data, maintaining the integrity of the dataset.
2. Categorical Encoding and Feature Engineering:

- Age Categorization: The age of individuals was converted into categorical variables to enhance model performance.
  - Occupation Simplification: The occupation feature, which had many unique values, was simplified to reduce complexity.
  - Blood Pressure Columns: The single blood pressure column was split into two separate columns for more detailed analysis.
  - BMI Categorization: The BMI index was converted from three categories into two for better model accuracy.
3. Data Visualization:
    - Appropriate visualizations were created to explore and understand the data patterns and relationships.
  4. Data Standardization and Labeling:
    - Data standardization was performed to normalize the features, and label encoding was applied to convert categorical variables into numerical format.
  5. Model Training and Evaluation:
    - An XGBoost (XGB) model was trained on the processed dataset, achieving an accuracy of 94.5%.

These steps ensured a comprehensive analysis and model training process, leading to a highly accurate prediction model for sleep disorders.

*Developed by Hosein Mohammadi*

GitHub : <https://github.com/Hosein541>

Kaggle : <https://www.kaggle.com/hoseinnnnnnn>

Gmail : [Huseinmohammadi83@gmail.com](mailto:Huseinmohammadi83@gmail.com)